

# AWS Data Lake and Analytics Platform

## Overview

I designed and implemented a data lake solution on AWS, enabling the organization to store, process, and analyze large volumes of structured and unstructured data.

## Key Components

- S3: For data storage
- Glue: For ETL jobs and data catalog
- Athena: For SQL queries on S3 data
- QuickSight: For data visualization
- Lambda: For data processing and transformations

## Implementation Highlights

### Glue ETL Job (PySpark)

python  
Copy

```
import sys
from awsglue.transforms import *
from awsglue.utils import getResolvedOptions
from pyspark.context import SparkContext
from awsglue.context import GlueContext
from awsglue.job import Job

glueContext = GlueContext(SparkContext.getOrCreate())
job = Job(glueContext)

# Read data from S3
datasource0 = glueContext.create_dynamic_frame.from_catalog(
    database = "my_data_lake_db",
    table_name = "raw_data"
)

# Apply transformations
applymapping1 = ApplyMapping.apply(
    frame = datasource0,
    mappings = [
        ("col1", "string", "transformed_col1", "string"),
        ("col2", "long", "transformed_col2", "int"),
        ("col3", "string", "transformed_col3", "string")
    ]
)

# Write transformed data back to S3
glueContext.write_dynamic_frame.from_options(
    frame = applymapping1,
```

```

connection_type = "s3",
connection_options = {"path": "s3://my-data-lake-bucket/transformed/"},
format = "parquet"
)

job.commit()

```

## Lambda Function for Data Processing

python  
Copy

```

import json
import boto3

s3 = boto3.client('s3')
athena = boto3.client('athena')

def lambda_handler(event, context):
    # Get the object from the event
    bucket = event['Records'][0]['s3']['bucket']['name']
    key = event['Records'][0]['s3']['object']['key']

    # Process the data (example: count words)
    response = s3.get_object(Bucket=bucket, Key=key)
    content = response['Body'].read().decode('utf-8')
    word_count = len(content.split())

    # Write results back to S3
    result = json.dumps({'file': key, 'word_count': word_count})
    s3.put_object(Bucket=bucket, Key=f'processed/{key}_count.json', Body=result)

    # Trigger Athena query
    query = f"ALTER TABLE processed_data ADD PARTITION (dt='{key.split('/')[0]}')"
    athena.start_query_execution(
        QueryString=query,
        QueryExecutionContext={'Database': 'my_data_lake_db'},
        ResultConfiguration={'OutputLocation': 's3://my-athena-results/'})

    return {
        'statusCode': 200,
        'body': json.dumps('Processing complete!')
    }

```

## Challenges and Solutions

- Data partitioning: Implemented a partitioning strategy in S3 for improved query performance.
- Cost optimization: Set up S3 lifecycle policies and used Athena workgroups to manage costs.

## **Outcome**

This project demonstrated my ability to architect and implement big data solutions on AWS, showcasing skills in data engineering and analytics.